
System Resource Utilization Monitoring for Docker Containers

By: Joseph Azevedo
and Bhanu Garg





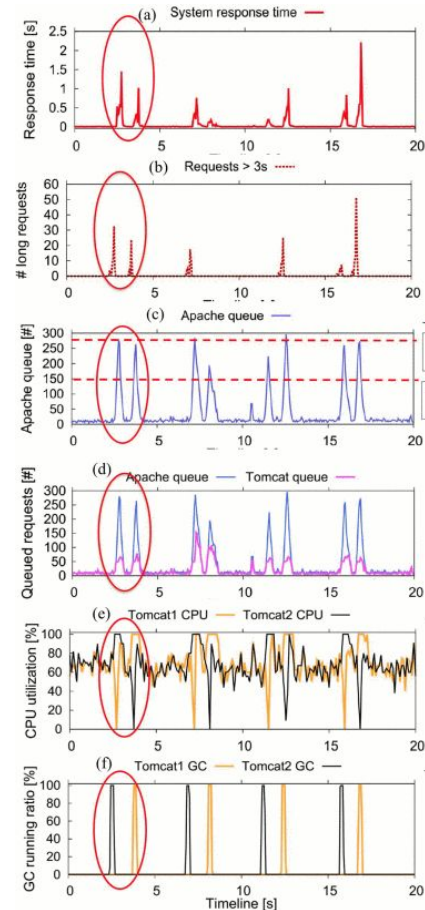
Motivation

- Increasing adoption of containerization technologies
 - 55% in 2017 to 87% in 2019
- Many solutions for monitoring containers
- Lack of tooling with **high granularity** and **low overhead**



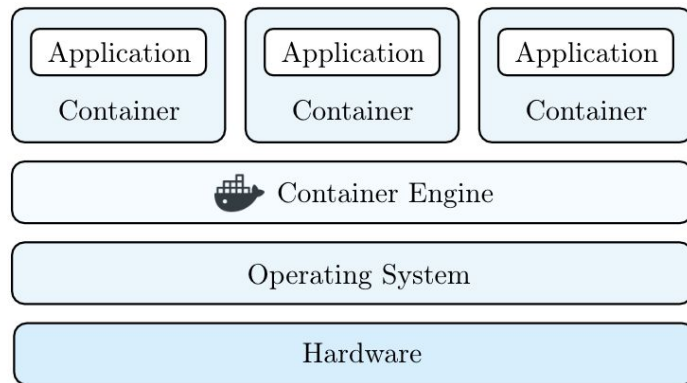
Millibottlenecks

- Brief resource bottlenecks that propagate through an application system
 - Short intervals
- Become amplified due to dependencies and their effect on downstream components
- Long tail problem
- Example: Google reported a loss in up to 20% of revenue from latencies > 500 milliseconds



Docker

- Package applications with their process-level configuration and dependencies
- Images can be instantiated as containers on a virtual environment
- Uses host OS's kernel






Project Evolution

- Original goal was to adapt the WISE microblog benchmark to run under Kubernetes and Docker
- Began to examine monitoring tools
 - Docker stats
 - Sysdig
 - cAdvisor
 - Moby



kubernetes



docker


Monitoring Tools

- **High Granularity**
 - consistent under load
 - precision of sampled data collected in each interval
- **Low Overhead**
 - low impact on application performance



Docker Stats

- Native to Docker
- Polls the Docker engine every 1 second
 - Can not be decreased
- Can not achieve high granularity





cAdvisor

- CLI tool developed by Google
- Geared toward system administrators
- Operates as a background daemon
 - Collects statistics about running containers
- Output log provides high granularity
 - Excessive monitoring leads to high overhead





cAdvisor

cAdvisor Output

cName=/system.slice/system-modprobe.slice host=localhost:8086

memory_usage=0 memory_working_set=0 rx_bytes=0 rx_errors=0

tx_bytes=0 tx_errors=0 **timestamp=1583873824045024464**

cpu_cumulative_usage=1435685

cName=/system.slice/systemd-logind.service host=localhost:8086

timestamp=1583873824056061533 **cpu_cumulative_usage=4974510763**

memory_usage=1798144 memory_working_set=1564672 rx_bytes=0

rx_errors=0 tx_bytes=0 tx_errors=0



Sysdig

- Curses based CLI tool
- Translating the outputs to a log file is difficult
- Collection interval changing does not work
 - Limited granularity



(C)Sysdig Output



```
(jazevedo) clnode055.clemson.cloudlab.us — Konsole

Viewing: Containers For: whole machine
Source: Live System Filter: container.name ≠ host
```

CPU	PROCS	THREADS	VIRT	RES	FILE	NET	ENGINE	IMAGE	ID	NAME
1244.00	17	17	4G	3G	0	0.00	docker	ubuntu	ce1d10669257	jovial_burnell

```
F1Help F2Views F4Filter F5Echo F6Dig F7Legend F8Actions F9Sort F12Spectro CTRL+FSearchp Pause 1/1(100.0%)
```

Moby

- Open-source framework that is part of the core Docker engine
 - Built in Golang
- Modified Moby fork to add configuration
 - Allowed to control collection granularity
- Drawbacks
 - Adds complexity, hard to extend



Implementation

rAdvisor





rAdvisor

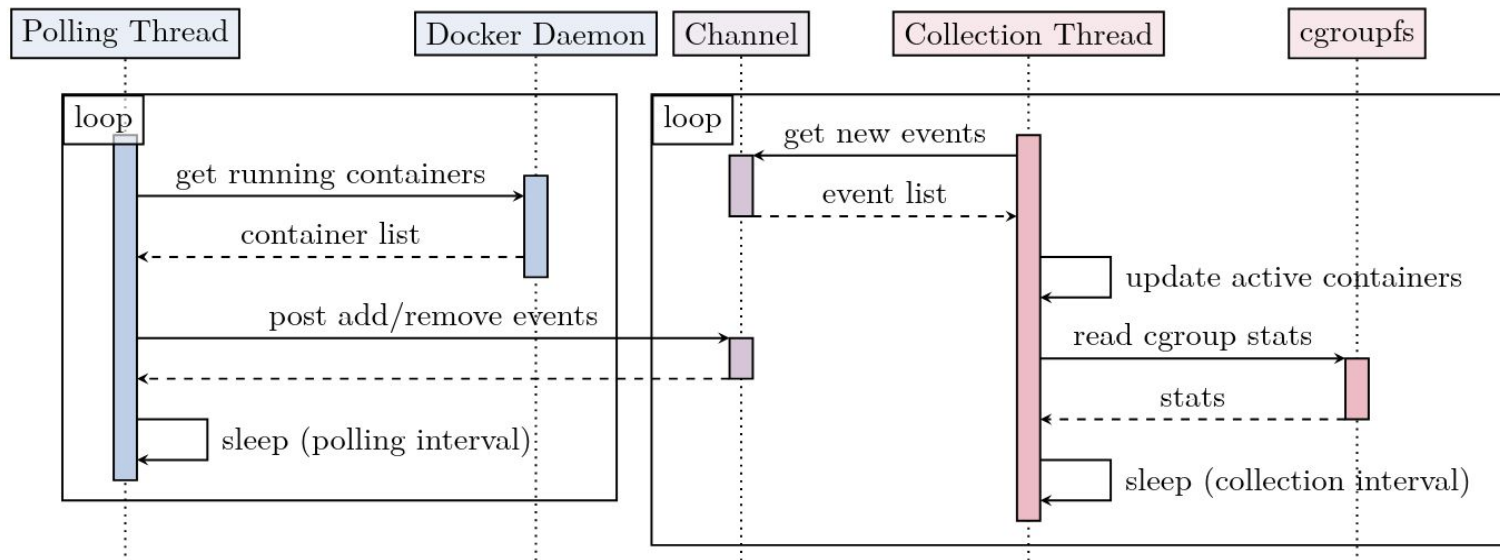
- Built in Rust
 - similar performance to C allows for lower overhead
 - no garbage collection allows for consistency in sampling
- Collects statistics for active containers at the specified collection interval
- Uses Linux cgroups to monitor containers
 - Operates efficiently to minimize overhead

rAdvisor



rAdvisor Structure

rAdvisor





Demo



Evaluation

rAdvisor vs Moby




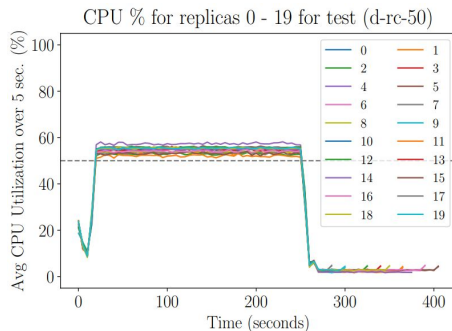
rAdvisor

- Primary criteria
 - Low overhead
 - High granularity
- Developed list of specific requirements to meet these
 - Low measurable performance impact when running
 - Consistent, low observed collection intervals



Experiments

- 3 different experiment configurations
 - 2 synthetic
 - 1 using WISE microblog benchmark
- Used experimental matrix
- Ran 30 replicas of each config



Tooling

- Experimental workflow scripts
- Automation scripts



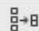
```
[ INFO 13:20:06.241] [ii-rc-s-24] Instantiating experiment ii-rc-s-24
(9fbbd8df-87c0-11ea-b1eb-e4434b2381fc) profile = MicroblogBareMetalC8220
[DEBUG 13:20:06.266] [ii-rc-s-24] Waiting for experiment to become ready
[ INFO 13:31:51.361] [ii-rc-s-24] Successfully provisioned new experiment from
cloudlab: MicroblogBareMetalC8220 (9fbbd8df-87c0-11ea-b1eb-e4434b2381fc) profile =
ii-rc-s-24
| clnode075.clemson.cloudlab.us...
[ INFO 13:31:51.362] [ii-rc-s-24] Using working/ii-rc-s-24 as the working directory
[ INFO 13:31:51.362] [ii-rc-s-24] Wrote rendered config file to
working/ii-rc-s-24/config.sh
```

Parsers and Analysis


- Python parsing scripts
- Jupyter Notebooks
 - Iterative analysis
 - Easy visualization



[9]

MI 

```
block_size = 20  
examine_single_replica_test("d-rc-50", replica_block_size=block_size, sampling_period=4.0, peak=200)
```





Results

Over 360 experiments ran

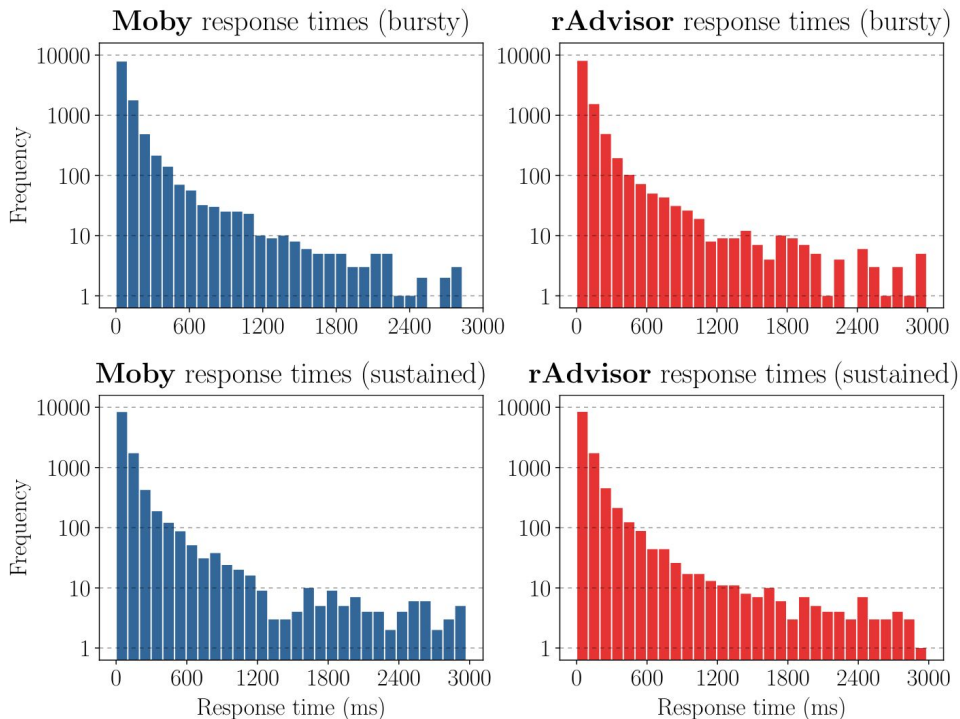
Synthetic Test Results

- Direct CPU utilization comparisons
 - rAdvisor had significantly less performance overhead than Moby under **load**
 - Moby has less performance overhead than rAdvisor at idle
 - This isn't as important
- Synthetic indirect test results
 - Moby had less of an indirect effect on CPU throughput in nbench benchmark

Aggregate Score	50 % peak CPU utilization		
	Control	Moby	rAdvisor
memory	29.16	-0.72 %	-2.03 %
integer	26.49	-0.77 %	-1.92 %
floating point	41.54	-0.74 %	-2.00 %



Aggregate Score	100 % peak CPU utilization		
	Control	Moby	rAdvisor
memory	29.17	-0.78 %	-2.09 %
integer	26.54	-0.77 %	-1.56 %
floating point	41.62	-0.81 %	-1.65 %



Response Time





- No noticeable difference between Moby and rAdvisor response times
- Observed emergence of the long tail phenomenon

Observed Collection Intervals










Future Work

- Extend/improve experimental workflow
 - Finish adapting WISE microblog benchmark to Kubernetes
 - Test additional configuration options, such as target collection interval
 - Add better synchronization measures in distributed experiments
- Extend/improve rAdvisor
 - Improve performance of polling thread
 - Use lightweight HTTP client or just get the list of active containers by listing the parent docker cgroup folder
 - Make it work in Windows with the Host Compute Services



kubernetes





Conclusion

- Investigated several monitoring tools
- Developed rAdvisor to detect Millibottlenecks
 - Achieved high granularity and low overhead
- Able to maintain at most 100ms collection intervals
 - Over 3 million collection samples
- Experienced low overhead during periods of load